

LECTURE 15: CURVE FITTING III

POLYNOMIAL REGRESSION
LOCAL FITTING



OCT 16 2025

INTERVAL “LINES”

- For any specific x_0 , the fitted (mean) value is:

- $$\hat{y}_0 = \hat{\beta}_0 + \hat{\beta}_1 x_0$$

- The standard error of that mean is:

- $$SE_{\hat{y}_0} = \sqrt{\sigma^2 \left(\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum (x_i - \bar{x})^2} \right)}$$

- The 95% confidence interval for the mean response is:


- $$\hat{y}_0 \pm t^* \times SE(\hat{y}_0)$$

- The lines for these intervals are **not parallel**


- They widen away from \bar{x} because $SE(\hat{y}_0)$ increases as x_0 moves away from the mean.
- More a “**bowtie**” shape

- And the prediction interval lines are bowties as well

- $$\hat{y}_0 \pm t^* \times SE(\text{pred}(x_0))$$



Salary vs Years of Experience




POLYNOMIAL REGRESSION

- We could a quadratic function:
 - $y = \beta_0 + \beta_1 x + \beta_2 x^2 + e$
- Determine optimal coefficients:
 - $\hat{y}_i(\beta_0, \beta_1, \beta_2) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2$,
- Error:
 - $\ell(y_i, \hat{y}_i(\beta_0, \beta_1, \beta_2))$
- Apply least squares
 - $\frac{1}{n} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i - \beta_2 x_i^2)^2$


SOLVING FOR POLYNOMIALS

- Polynomial regression is still *linear in parameters*: we just include X^2 , X^3 , ... as extra columns in X .
- The goal is the same: find coefficients β_0 , β_1 , β_2 , ... that minimize the sum of squared errors.
- Matrix form: $Y = X\beta + \varepsilon$ where $X = [1, X, X^2, \dots]$.
- Least-squares solution: $\hat{\beta} = (X^T X)^{-1} X^T Y$.
- So polynomial regression simply solves a **larger linear system** using the same principle as simple linear regression.


Salary vs Years of Experience



Salary vs Years of Experience



Salary vs Years of Experience



INTERPRETATION OF POLYNOMIAL TERMS

- β_1, β_2, \dots describe curvature, not direct linear effects
- Each higher-order term refines fit near extremes of X
- Interpretation focuses on overall curve shape, not single coefficient
- Higher degrees increase flexibility but risk overfitting
- Choose degree using residual plots or cross-validation
- Underfitting: trend missed; Overfitting: noise captured

POLYNOMIAL REGRESSION

- In polynomial regression, the design matrix X has columns $[1, x, x^2, x^3, \dots]$
- **Model setup**
 - Polynomial regression model: $Y = X\beta + \varepsilon$, where $X = [1, x, x^2, \dots, x^d]$.
 - The goal is to find the coefficients β that minimize total squared error.
- **Define the loss function**
 - We measure total squared residuals using $L(\beta) = (Y - X\beta)^T (Y - X\beta)$.
 - This is like adding up all squared differences between the observed and predicted values.
- **Expand the expression**
 - When expanded, the loss becomes $L = Y^T Y - 2\beta^T X^T Y + \beta^T X^T X\beta$.
 - Each term reflects a different interaction between data, model, and parameters

SOLVING

- Differentiate the loss
 - Using matrix calculus rules:
 - If $L = b^\top \beta$, then its derivative is b
 - If $L = \beta^\top A \beta$ where A is symmetric, the derivative is $2A\beta$
- Simplify and solve
 - The derivative of the loss is $\partial L / \partial \beta = -2X^\top Y + 2X^\top X\beta$
 - Setting this equal to zero gives the minimum: $X^\top X\beta = X^\top Y$
 - Solving for β gives $\hat{\beta} = (X^\top X)^{-1} X^\top Y$
- Key idea
 - Even though the predictors include X, X^2, X^3 and so on, the model is still linear in β
 - So the least-squares solution works exactly like in simple linear regression, only in higher dimensions

LINEAR ALGEBRA

■ Transpose (X^T)

- The transpose flips rows and columns so matrix multiplication can work correctly.
- Example: $x = [[1, 2, 3], [4, 5, 6]]$ becomes $x^T = [[1, 4], [2, 5], [3, 6]]$
- Like rotating the data table so that all variables align for dot products.

■ Inverse (A^{-1})

- The inverse is the matrix version of dividing by a number.
- For numbers, dividing by 4 is the same as multiplying by $\frac{1}{4}$.
- For matrices, multiplying by A^{-1} undoes the effect of multiplying by A, since $A^{-1}A = I$.
- In regression, $(X^T X)^{-1}$ removes overlap among predictors, so each coefficient β is adjusted for the others.
It is like dividing out the entanglement among predictors.


■ Pseudoinverse (X^+)

- Defined as $(X^T X)^{-1} X^T$ when $X^T X$ can be inverted.
- This is the matrix equivalent of “smart division” by X, used when X is not square or directly invertible.
- It gives the compact solution $\hat{\beta} = X^+ Y$

LOCAL FITTING

FUNCTIONS

- What would be a good fit here ?




LOCAL FITTING OVERVIEW

- We can have a **single complex** function
 - Or, combine **multiple simple** functions
- Local fitting methods estimate the relationship between X and Y using **nearby data points**
 - Rather than a single global model
- Each fitted value is based on a small neighborhood around x_0 .
- This produces flexible, smooth curves that adapt to local structure in the data

SLIDING WINDOWS RUN THE MEAN (OR MEDIAN)

- $\hat{f}(x) = \frac{1}{\#\text{in window}} \sum_{i: x_i \in [x - \frac{w}{2}, x + \frac{w}{2})} y_i$




KERNEL WEIGHTING


$$\hat{f}(x) = \frac{\sum_{i:x_i \in [x - \frac{w}{2}, x + \frac{w}{2}]} y_i}{\sum_{i:x_i \in [x - \frac{w}{2}, x + \frac{w}{2}]} 1}$$

- $$= \frac{\sum_{i=1}^n y_i f(x, x_i)}{\sum_{i=1}^n f(x, x_i)}$$
- $$f(x, x_i) = \begin{cases} \frac{1}{w} & x_i \in [x - \frac{w}{2}, x + \frac{w}{2}] \\ 0 & otherwise \end{cases}$$

- Nadarya-Watson Kernel
- How does it compare with running mean/ median ?
- What is the effect of the window size ?




GAUSSIAN KERNEL




- Is a popular choice

Moving Average



Kernel Smoothing



MATHEMATICAL FORMULATION

- At **each** target x_0
 - Estimate coefficients $\beta_0(x_0)$, $\beta_1(x_0)$ by minimizing
 - $\sum K((x_i - x_0)/h) \times (y_i - \beta_0 - \beta_1 x_i)^2$
- K is a kernel function **assigning higher weight to nearby x_i**
- h (bandwidth) controls how wide the neighborhood is


KERNEL WEIGHTING FUNCTIONS


- Kernel weights determine how quickly influence declines with distance from x_0 .
- Common choices: Gaussian, Epanechnikov, Tricube kernels
- Larger $|x_i - x_0| \Rightarrow$ smaller weight $K((x_i - x_0)/h)$





LOESS









Zoomed up




Zoomed up




Mean



Linear Regression



Quadratic Regression




LOESS

- LOESS (Locally Estimated Scatterplot Smoothing) fits **many small regressions** instead of one global line
- At each target point x_0 , it fits a **local polynomial** (usually linear or quadratic) using **weighted nearby points**
- Nearby points get **more weight**, distant points get **less weight**
- The fitted value at x_0 is the prediction from that local weighted regression
- It is a **non-parametric** method, we don't assume one global β for the entire dataset

THE MATH

- At each x_0
 - We minimize the **weighted least-squares loss**:
 - $L(\beta \mid x_0) = \sum_i w_i(x_0) [y_i - (\beta_0 + \beta_1 x_i + \beta_2 x_i^2 \dots)]^2$
 - The weight $w_i(x_0)$ is **largest near x_0**
 - The **Tricube** kernel
 - $w_i(x_0) = (1 - |x_i - x_0|^3 / d^3)^3 \quad \text{for } |x_i - x_0| < d, \text{ else } 0$
- The solution at x_0 gives the local coefficients:
$$\hat{\beta}(x_0) = (X^T W(x_0) X)^{-1} X^T W(x_0) Y$$
- The fitted value:
$$\hat{y}(x_0) = [1, x_0, x_0^2, \dots] \hat{\beta}(x_0)$$



SOLUTION

- We solve: $\hat{\beta}(x_0) = (X^T W X)^{-1} X^T W Y$

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ 1 & x_3 \end{bmatrix}, \quad W = \begin{bmatrix} w_1 & 0 & 0 \\ 0 & w_2 & 0 \\ 0 & 0 & w_3 \end{bmatrix}, \quad Y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix}$$

Then:

$$\hat{\beta}(x_0) = (X^T W X)^{-1} X^T W Y$$

This yields local slope and intercept (β_0, β_1) at point x_0 .

EXAMPLE

- Suppose we have 3 nearby points around $\mathbf{x}_0 = 5$
 - $\mathbf{x} = [4, 5, 6]$
 - $\mathbf{y} = [10, 15, 18]$
- Weights using the tricube kernel):
 - $\mathbf{w} = [0.25, 1.0, 0.25]$
- Then $\mathbf{x}_0 = 5$ gets the most weight !

$$X = \begin{bmatrix} 1 & 4 \\ 1 & 5 \\ 1 & 6 \end{bmatrix}, \quad W = \begin{bmatrix} 0.25 & 0 & 0 \\ 0 & 1.0 & 0 \\ 0 & 0 & 0.25 \end{bmatrix}$$

$$\hat{\beta} = (X^T W X)^{-1} X^T W Y$$

$$\hat{\beta} \approx \begin{bmatrix} -5.0 \\ 3.9 \end{bmatrix}$$

$$\hat{y}(5) = -5.0 + 3.9 \times 5 = 14.5$$

LOESS SUMMARY

- LOESS repeats a local weighted regression at each x_0
- Each local fit has its own $\beta_0(x_0)$, $\beta_1(x_0)$, etc.
- Nearby points dominate each local line; faraway points barely influence it
- The collection of all local predictions forms a **smooth curve** that adapts to the data
 - The final “curve” is a combination of small/short straight lines (chain links)